

# Mercedes-Benz Greener Manufacturing Course-end Project 1

October 16, 2023

## 1 Mercedes-Benz Greener Manufacturing

Course-end Project 1

Name-Nishant Dubey

Objective-You are required to reduce the time that cars spend on the test bench. Others will work with a dataset representing different permutations of features in a Mercedes-Benz car to predict the time it takes to pass testing. Optimal algorithms will contribute to faster testing, resulting in lower carbon dioxide emissions without reducing Mercedes-Benz's standards

```
[1]: # Importing the required libraries

import numpy as np
import pandas as pd
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: # Importing the data

train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')
```

```
[3]: train.head()
```

```
[3]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	\
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	

  

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0

```

2      0      1      0      0      0
3      0      0      0      0      0
4      0      0      0      0      0

```

[5 rows x 378 columns]

```
[4]: test.head()
```

```

[4]:   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10  ...  X375  X376  X377  X378  X379  X380  \
0    1  az  v   n  f  d  t  a  w    0  ...    0    0    0    1    0    0
1    2   t  b  ai  a  d  b  g  y    0  ...    0    0    1    0    0    0
2    3  az  v  as  f  d  a  j  j    0  ...    0    0    0    1    0    0
3    4  az  l   n  f  d  z  l  n    0  ...    0    0    0    1    0    0
4    5   w  s  as  c  d  y  i  m    0  ...    1    0    0    0    0    0

```

```

      X382  X383  X384  X385
0         0     0     0     0
1         0     0     0     0
2         0     0     0     0
3         0     0     0     0
4         0     0     0     0

```

[5 rows x 377 columns]

```
[5]: train.columns
```

```

[5]: Index(['ID', 'y', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8',
...
      'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
      'X385'],
      dtype='object', length=378)

```

```
[6]: test.columns
```

```

[6]: Index(['ID', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8', 'X10',
...
      'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
      'X385'],
      dtype='object', length=377)

```

```

[7]: print('Size of training set: {} rows and {} columns'.format(*train.shape))
      print('Size of testing set: {} rows and {} columns'.format(*test.shape))

```

```

Size of training set: 4209 rows and 378 columns
Size of testing set: 4209 rows and 377 columns

```

```

[8]: # Collect the Y values into an array
      y_train = train['y'].values

```

```
[9]: y_train
```

```
[9]: array([130.81,  88.53,  76.26, ..., 109.22,  87.48, 110.85])
```

```
[10]: # Understanding the data types:
cols = [c for c in train.columns if 'X' in c]
print('Number of features: {}'.format(len(cols)))
print('Feature types:')
train[cols].dtypes.value_counts()
```

Number of features: 376

Feature types:

```
[10]: int64    368
      object     8
      dtype: int64
```

```
[11]: # Count the data in each of the columns

counts = [[], [], []]
for c in cols:
    typ = train[c].dtype
    uniq = len(np.unique(train[c]))
    if uniq == 1:
        counts[0].append(c)
    elif uniq == 2 and typ == np.int64:
        counts[1].append(c)
    else:
        counts[2].append(c)
print('Constant features: {} Binary feature: {} Categorical features: {}'.format(*[len(c) for c in counts]))
print('Constant features:', counts[0])
print('Categorical features:', counts[2])
```

Constant features: 12 Binary feature: 356 Categorical features: 8

Constant features: ['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330', 'X347']

Categorical features: ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']

```
[12]: # Splitting the data

usable_columns = list(set(train.columns) - set(['ID', 'y']))
y_train = train['y'].values
id_test = test['ID'].values
x_train = train[usable_columns]
x_test = test[usable_columns]
```

## 2 Checking for null values and unique values for train and test data

```
[13]: def check_missing_values(df):  
      if df.isnull().any().any():  
          print('There are missing values in the dataframe')  
      else:  
          print('There are no missing values in the dataframe')
```

```
[14]: check_missing_values(x_train)  
      check_missing_values(x_test)
```

There are no missing values in the dataframe  
There are no missing values in the dataframe

## 3 Label Encoding the categorical values

```
[15]: for column in usable_columns:  
      cardinality = len(np.unique(x_train[column]))  
      if cardinality == 1:  
          x_train.drop(column, axis=1) # column with only one  
          # value is useless so we drop it.  
          x_test.drop(column, axis=1)  
      if cardinality > 2: # Column is categorical  
          mapper = lambda x: sum([ord(digit) for digit in x])  
          x_train[column] = x_train[column].apply(mapper)  
          x_test[column] = x_test[column].apply(mapper)  
x_train.head()
```

```
[15]:
```

	X24	X356	X291	X274	X292	X324	X178	X325	X142	X308	...	X196	X107	\
0	0	0	0	0	0	1	0	0	1	0	...	0	0	
1	0	0	0	0	0	0	1	0	1	0	...	0	0	
2	0	0	0	1	0	1	0	0	0	0	...	0	0	
3	0	0	1	0	0	0	0	0	1	0	...	0	0	
4	0	0	0	0	0	0	0	0	0	0	...	0	0	

  

	X385	X187	X311	X212	X217	X164	X355	X314
0	0	1	0	0	0	0	0	0
1	0	1	1	0	0	0	0	0
2	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0

[5 rows x 376 columns]

```
[16]: # Make sure the data is changed into numerical values
```

```
print('featurectypes:')
x_train[cols].dtypes.value_counts()
```

featurectypes:

```
[16]: int64    376
      dtype: int64
```

## 4 Perform Dimensionality Reduction

```
[17]: n_comp = 12
      pca = PCA(n_components = n_comp, random_state = 420)
      pca2_results_train = pca.fit_transform(x_train)
      pca2_results_test = pca.transform(x_test)
```

## 5 Training Using XGBoost

```
[18]: # Training Using XGBoost

import xgboost as xgb
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split
```

```
[19]: x_train, x_val, y_train, y_val = train_test_split(pca2_results_train, y_train,
      ↪ test_size=0.2, random_state=4242)
```

```
[20]: d_train = xgb.DMatrix(x_train, label = y_train)
      d_val = xgb.DMatrix(x_val, label = y_val)

      # d_test = xgb.DMatrix(x_test)

      d_test = xgb.DMatrix(pca2_results_test)
```

```
[21]: params = {}
      params['Objective'] = 'reg:linear'
      params['eta'] = 0.02
      params['max_depth'] = 4

      def xgb_r2_score(preds, dtrain):
          labels = dtrain.get_label()
          return 'r2', r2_score(labels, preds)
      watchlist = [(d_train, 'train'), (d_val, 'valid')]
      clf = xgb.train(params, d_train, 1000, watchlist, early_stopping_rounds=50,
          feval=xgb_r2_score, maximize=True, verbose_eval=10)
```

[13:09:23] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0fdc6d574b9c0d168-1\xgboost\xgboost-ci-windows\src\learner.cc:767: Parameters: { "Objective" } are not used.

[0]	train-rmse:99.14834	train-r2:-58.35295	valid-rmse:98.26297
	valid-r2:-67.63754		
[10]	train-rmse:81.27653	train-r2:-38.88428	valid-rmse:80.36433
	valid-r2:-44.91014		
[20]	train-rmse:66.71610	train-r2:-25.87403	valid-rmse:65.77334
	valid-r2:-29.75260		
[30]	train-rmse:54.86913	train-r2:-17.17722	valid-rmse:53.89147
	valid-r2:-19.64534		
[40]	train-rmse:45.24710	train-r2:-11.36098	valid-rmse:44.22334
	valid-r2:-12.90225		
[50]	train-rmse:37.44856	train-r2:-7.46723	valid-rmse:36.37638
	valid-r2:-8.40634		
[60]	train-rmse:31.14585	train-r2:-4.85695	valid-rmse:30.02279
	valid-r2:-5.40743		
[70]	train-rmse:26.08417	train-r2:-3.10795	valid-rmse:24.91516
	valid-r2:-3.41275		
[80]	train-rmse:22.04312	train-r2:-1.93371	valid-rmse:20.83299
	valid-r2:-2.08521		
[90]	train-rmse:18.84671	train-r2:-1.14458	valid-rmse:17.59846
	valid-r2:-1.20156		
[100]	train-rmse:16.33186	train-r2:-0.61043	valid-rmse:15.08617
	valid-r2:-0.61786		
[110]	train-rmse:14.39874	train-r2:-0.25176	valid-rmse:13.15521
	valid-r2:-0.23020		
[120]	train-rmse:12.92910	train-r2:-0.00927	valid-rmse:11.70051
	valid-r2:0.02682		
[130]	train-rmse:11.81536	train-r2:0.15712	valid-rmse:10.62244
	valid-r2:0.19790		
[140]	train-rmse:10.99099	train-r2:0.27063	valid-rmse:9.86019
	valid-r2:0.30888		
[150]	train-rmse:10.38667	train-r2:0.34863	valid-rmse:9.33123
	valid-r2:0.38104		
[160]	train-rmse:9.93418	train-r2:0.40415	valid-rmse:8.96192
	valid-r2:0.42907		
[170]	train-rmse:9.59640	train-r2:0.44398	valid-rmse:8.71810
	valid-r2:0.45971		
[180]	train-rmse:9.35220	train-r2:0.47192	valid-rmse:8.55750
	valid-r2:0.47943		
[190]	train-rmse:9.16592	train-r2:0.49275	valid-rmse:8.45262
	valid-r2:0.49212		
[200]	train-rmse:9.02357	train-r2:0.50838	valid-rmse:8.38960
	valid-r2:0.49966		
[210]	train-rmse:8.92419	train-r2:0.51915	valid-rmse:8.35118
	valid-r2:0.50423		

[220]	train-rmse:8.84149 valid-r2:0.50685	train-r2:0.52802	valid-rmse:8.32911
[230]	train-rmse:8.77383 valid-r2:0.50892	train-r2:0.53522	valid-rmse:8.31164
[240]	train-rmse:8.72642 valid-r2:0.51010	train-r2:0.54023	valid-rmse:8.30160
[250]	train-rmse:8.68650 valid-r2:0.51034	train-r2:0.54442	valid-rmse:8.29958
[260]	train-rmse:8.64705 valid-r2:0.51107	train-r2:0.54855	valid-rmse:8.29340
[270]	train-rmse:8.61922 valid-r2:0.51093	train-r2:0.55145	valid-rmse:8.29457
[280]	train-rmse:8.58611 valid-r2:0.51118	train-r2:0.55489	valid-rmse:8.29251
[290]	train-rmse:8.55652 valid-r2:0.51121	train-r2:0.55796	valid-rmse:8.29217
[300]	train-rmse:8.53319 valid-r2:0.51117	train-r2:0.56036	valid-rmse:8.29254
[310]	train-rmse:8.50784 valid-r2:0.51129	train-r2:0.56297	valid-rmse:8.29156
[320]	train-rmse:8.48199 valid-r2:0.51135	train-r2:0.56562	valid-rmse:8.29100
[330]	train-rmse:8.45003 valid-r2:0.51155	train-r2:0.56889	valid-rmse:8.28928
[340]	train-rmse:8.42263 valid-r2:0.51171	train-r2:0.57168	valid-rmse:8.28797
[350]	train-rmse:8.39358 valid-r2:0.51183	train-r2:0.57463	valid-rmse:8.28693
[360]	train-rmse:8.37163 valid-r2:0.51188	train-r2:0.57685	valid-rmse:8.28655
[370]	train-rmse:8.34326 valid-r2:0.51201	train-r2:0.57972	valid-rmse:8.28542
[380]	train-rmse:8.31805 valid-r2:0.51219	train-r2:0.58225	valid-rmse:8.28393
[390]	train-rmse:8.28994 valid-r2:0.51239	train-r2:0.58507	valid-rmse:8.28216
[400]	train-rmse:8.26600 valid-r2:0.51254	train-r2:0.58746	valid-rmse:8.28089
[410]	train-rmse:8.24563 valid-r2:0.51247	train-r2:0.58949	valid-rmse:8.28154
[420]	train-rmse:8.21978 valid-r2:0.51250	train-r2:0.59206	valid-rmse:8.28123
[430]	train-rmse:8.19649 valid-r2:0.51268	train-r2:0.59437	valid-rmse:8.27975
[440]	train-rmse:8.17680 valid-r2:0.51260	train-r2:0.59632	valid-rmse:8.28042
[450]	train-rmse:8.15583 valid-r2:0.51265	train-r2:0.59839	valid-rmse:8.27997

```

[460]   train-rmse:8.13170      train-r2:0.60076      valid-rmse:8.27872
valid-r2:0.51280
[470]   train-rmse:8.10759      train-r2:0.60312      valid-rmse:8.28004
valid-r2:0.51264
[480]   train-rmse:8.08873      train-r2:0.60497      valid-rmse:8.27962
valid-r2:0.51269
[490]   train-rmse:8.06167      train-r2:0.60761      valid-rmse:8.28014
valid-r2:0.51263
[500]   train-rmse:8.03613      train-r2:0.61009      valid-rmse:8.27783
valid-r2:0.51290
[510]   train-rmse:8.01630      train-r2:0.61201      valid-rmse:8.27981
valid-r2:0.51267
[520]   train-rmse:7.98437      train-r2:0.61510      valid-rmse:8.28011
valid-r2:0.51263
[530]   train-rmse:7.96313      train-r2:0.61714      valid-rmse:8.28034
valid-r2:0.51261
[540]   train-rmse:7.93430      train-r2:0.61991      valid-rmse:8.28030
valid-r2:0.51261
[550]   train-rmse:7.91141      train-r2:0.62210      valid-rmse:8.28270
valid-r2:0.51233
[554]   train-rmse:7.90739      train-r2:0.62248      valid-rmse:8.28240
valid-r2:0.51237

```

## 6 Predicting test\_df using XGBoost

```
[22]: p_test = clf.predict(d_test)
```

```
[23]: sub = pd.DataFrame()
sub['ID'] = id_test
sub['y'] = p_test
sub.to_csv('test_df.csv', index = False)
sub.head()
```

```
[23]:
```

	ID	y
0	1	83.397812
1	2	97.286064
2	3	83.171097
3	4	76.930611
4	5	112.544647

```
[ ]:
```